



MITCHELL WEISS

WILLIAM MARKS

NVIDIA and the Sovereignty Claws

Three days before NVIDIA GTC 2026, the tech giant's annual developer conference in March, and a highly anticipated address by NVIDIA's CEO, Jensen Huang, the company previewed both by posting on the social media platform X, "Get ready for a CLAW-some NVIDIA GTC."¹ The post made it almost certain that Huang would address the rise of AI "claws," a new category of more autonomous AI agents, and look ahead to the nature of NVIDIA's engagement in the space, whatever it would be.

The rapidly evolving AI agent scene invited Huang's views on claws, but so did a shifting landscape on AI "sovereignty." Control over components of AI – data, models, inference, etc. – was already of great interest among countries and companies. New national protection provisions, which had hardly been tested by the current wave of generative AI, would be further pressed by the rise of AI agents and now "claws" which connected more, shared more, and acted more. So would corporate anxiety.

A month earlier, another large figure, India's Prime Minister Narendra Modi, had taken the stage at his own conference, the India AI Impact Summit 2026, to expound his vision for AI, with implications for NVIDIA and other AI giants. He paired an expansive vision for AI's transformative effects and India's ambitions helping to bring them about with comments about the need for direction, control, and standards.² Modi didn't make explicit reference to a key clause in India's Digital Personal Data Protection (DPDP) Act in his keynote remarks, but he didn't have to. Major AI players knew that Phase 2 of the Act was expected to take effect in November 2026 and would likely put many large tech companies on the hook for compliance with national provisions around data consent and data sharing as it related to agents. Though given the pace of change, when passed, the Act hadn't anticipated them.

A week ahead of Huang's talk at GTC, word was out that NVIDIA was planning a new platform called NemoClaw to make the most of the arriving claws, in part by creating a framework for control and security. How NemoClaw might address requirements of India's data privacy initiatives, let alone the country's AI sovereignty work, remained to be seen by Indian officials. The same was true for leaders in dozens of other countries also with sovereign AI efforts and their corporate counterparts.

Signs that NVIDIA was embracing claws, and the most famous one so far, the open-source OpenClaw, also suggested strategic riddles to parse. The company had climbed to a \$5 trillion valuation, in large part by selling its GPUs in giant numbers especially to a handful of giant tech companies. NVIDIA's emphasis on its chips and its proprietary software had helped secure almost unimaginable profits in recent years. Should Huang, with NemoClaw, signal a new direction? As the

Professor Mitchell Weiss and Lecturer William Marks prepared this case. This case was developed from published sources. Funding for the development of this case was provided by Harvard Business School and not by the company. HBS cases are developed solely as the basis for class discussion. Cases are not intended to serve as endorsements, sources of primary data, or illustrations of effective or ineffective management.

Copyright © 2026 President and Fellows of Harvard College. To order copies or request permission to reproduce materials, call 1-800-545-7685, write Harvard Business School Publishing, Boston, MA 02163, or go to www.hbsp.harvard.edu. This publication may not be digitized, photocopied, or otherwise reproduced, posted, or transmitted, without the permission of Harvard Business School.

proliferation of AI agents brought about increased concerns about AI sovereignty, but also opportunity, it seemed fair to ask about GTC, just how CLAW-some would it be?

NVIDIA

At the outset of March 2026, NVIDIA was the most valuable company in the world. It was then valued at close to \$4.5 trillion. (It had become the first to cross the \$5 trillion mark in October 2025 before markets declined.³) (See **Exhibit 1** for NVIDIA stock history.) The company had finished its fiscal 2026 with \$216 billion in revenue, up 65% from fiscal 2025.⁴ (See **Exhibit 2** for NVIDIA's income statement). It had earned \$43 billion in net income and returned \$41 billion to shareholders via repurchased shares and dividends.⁵ In topping all companies in market valuation, it had surpassed major customers – Alphabet, Microsoft, Amazon, Meta, Tesla – and its main supplier Taiwan Semiconductor Manufacturing Company (TSMC).

The company's roots were not in AI, but in gaming. NVIDIA was founded in 1993 by Huang, Chris Malachowsky, and Curtis Priem.⁶ "Chris and Curtis were building chips at Sun Microsystems, and I was the assigned engineer to work with them from LSI Logic," Huang recalled.⁷ "Chris and Curtis reached out to me and asked me if I would like to leave LSI Logic and join them to build a company," he explained. "We were excited about the PC revolution, and we thought, 'What application would we bring to the PC?' Well, the one thing you would do with it more than anything in the world is play games."⁸ Huang, who would go on to become the company's President and CEO recounted, "The capability, the graphics capability, the multi-media capability of the PC was non-existent at the time. We thought, 'Maybe 3D graphics would be the thing that'd be really cool.'" He added, "We just need[ed] to go build a chip that makes it possible to play."⁹ NVIDIA would eventually release the GeForce 256 in 1999 and introduce it "as the world's first GPU."¹⁰ It would handle the calculations for making the 3D world of games look accurate, taking care of how light bounced off surfaces, how shadows fell, etc. The CPU (central processing unit) in a gamer's personal computer would still handle running the game itself, its logic, scoring, etc. NVIDIA would go on to dominate this space with successive generations of GPUs (graphics processing units) and other graphic-related innovations. By 2022 their graphics segment was generating \$16 billion in annual sales.¹¹

The GPU's usefulness turned out to extend well beyond gaming. As the company described it, "The parallel power that transformed gaming graphics caught the attention of researchers, who realized these GPUs could also unlock massive computational potential in AI, enabling breakthroughs far beyond the gaming world."¹² (See **Exhibit 3a and 3b** for NVIDIA's architecture timeline and current Blackwell product lineup.) Because GPUs could undertake many operations at the same time, they proved very useful for training neural networks that became so foundational to AI advancements in the 2000s. Moreover, NVIDIA's CUDA software and programming framework, launched in 2006, made it substantially easier to program GPUs and to scale to multiple GPUs inside one machine.¹³ (A sizable AI model was too big for one GPU. And in any event, training on a single GPU would take years.) (See **Exhibit 4** for CUDA overview.) Other NVIDIA innovations meant hundreds and then hundreds of thousands of GPUs could be joined in clusters and work in parallel.

OpenAI's launch of ChatGPT in 2022 in many ways set off a race for access to NVIDIA's GPUs that had not abated. At least two factors were driving demand: major model providers continued to train new frontier models, and each new model required vast GPU clusters running for months. Furthermore, AI consumers were using the AI models for more inference, and this too leveraged GPUs. Not only were AI users growing in number, but new techniques were adding to the demand for compute. Huang had said at the 2025 GTC that new reasoning models and coming AI agents would

drive demand for inference more than 100-fold.¹⁴ Some analysts estimated inference demands would be close to 120 times training demands by 2026.¹⁵ Amazon, Alphabet, Meta, Microsoft, and OpenAI were projected “to spend more than \$325 billion on giant data centers” in 2025 alone.¹⁶ Elon Musk’s xAI had built a supercomputer, “Colossus,” in Memphis, Tennessee, that started with 100,000 NVIDIA H100 GPUs and then doubled to 200,000. They were said to be planning for one million there.¹⁷ Colossus and data centers like it had so grown in number and footprint – real estate and energy – that some communities had begun pushing back against their development.¹⁸ Demand from the companies seemed more definitive. Sam Altman, OpenAI’s CEO said in September 2025, “you should expect OpenAI to spend trillions of dollars on things like data center construction in the not-too-distant future.”¹⁹

At least for the time being, NVIDIA had dominated this market. It was said to hold a more than 90% share of the AI chip market, with Intel and AMD splitting most of the rest.²⁰ NVIDIA’s H100 GPU, launched in 2022 and sold for between \$24,000 and \$40,000 depending on configurations, had been the industry standard. The company’s Blackwell family of GPUs, introduced in 2024, had since become the chips of choice for companies training foundation models. On introduction, it was said to enable the building and running of trillion-parameter models “at up to 25x less cost and energy consumption” than its predecessor.²¹ The leaders of Google DeepMind, Meta, Microsoft, OpenAI, Oracle, Dell, and xAI were all quoted in its launch announcement. Musk had expressed his views most succinctly. “There is currently nothing better than NVIDIA hardware for AI.”²² For how long that would remain was an open question. OpenAI, Meta, Amazon, Anthropic, and others were each, either on their own or in partnership with Broadcom or others, developing their own AI chips.²³ NVIDIA acknowledged the risk: “For fiscal 2026, sales to one direct customer represented 22% of total revenue and sales to another direct customer represented 14% of total revenue.”²⁴ One outside article captured some of the mood, “Nvidia’s Big Tech customers might also be its biggest competitive threat.”

Three decades after its founding, NVIDIA was a much-changed company. Its graphics business produced \$22 billion in revenue. Its computer and networking business, which included the revenue from data center computing, stood at \$130 billion.²⁵ Less than a month before Huang would give his keynote at the 2026 GTC, the company had predicted \$78 billion in revenue for the next quarter alone, with \$73 billion projected to come from its data center business.²⁶

India’s National AI Plan

Just weeks before Huang would take the stage at his GTC, Modi took the stage at his India AI Impact Summit and articulated the duality of India’s AI approach: broad ambition and sovereign guardrails.

Modi laid out an expansive vision for the potential for AI in India and beyond. “AI is making machines intelligent, but more importantly, it is multiplying human capability many times over.”²⁷ Modi went on, “India does not see fear in AI. India sees fortune in AI. India sees the future in AI.”

The Prime Minister appeared to ground that ambition, however, in a policy lens. “We must give AI the open sky, but keep the command in our own hands,” he announced. He then presented what he called India’s vision for AI: M.A.N.A.V. (Manav meant “human” in Hindi.):

“The MANAV vision stands for:

M – Moral and Ethical Systems: means AI must be built on ethical guidelines.

A – Accountable Governance: means Transparent rules and robust oversight.

N – National Sovereignty: means Data belongs to those who generate it.

A – Accessible and Inclusive: means AI must not be a monopoly, but a multiplier.

V –Valid and Legitimate: means AI must be lawful and verifiable.”²⁸

Modi said, “India’s MANAV Vision will be a vital link for human welfare in the AI-driven world of the 21st century.”²⁹ (See **Exhibit 5** for Modi’s M.A.N.A.V. communication.)

Towards Data Sovereignty

India’s focus on data sovereignty was not new. In 2011, the IT Reasonable Security Practices and Procedures and Sensitive Personal Data or Information (SPDI) Rules introduced categories of sensitive personal data and required companies to obtain permission before collecting any of them.³⁰ In 2017, a constitutional bench of the Supreme Court of India unanimously ruled that the right to privacy was a fundamental right enshrined in Article 21 of the Constitution – the right to life and personal liberty.³¹

Modi’s government produced a draft of a Digital Personal Data Protection Bill (DPDP) in November 2022 for public consultation.³² It was shorter, more principles-based, and stripped of some complexity of predecessor versions. The government abandoned strict data localization in favor of a “blacklist” model: personal data could flow to any country outside India *except* those the central government specifically restricted by notification.³³ This shift reflected the government’s desire to remain globally competitive for technology investment while retaining sovereign discretion over data flows.

The DPDP Act included three phases, with Phase 1 effective November 2025, Phase 2 effective November 2026 and Phase 3 effective November 2027.³⁴ Phase 1 was primarily procedural, formally establishing the Data Protection Board of India (DPBI) as the independent adjudicator at center of DPDP. Phase 2 would mark the operational beginning of enforcement for data protections with the introduction of local Consent Managers required to allow individuals to manage data consents across multiple platforms. Finally, Phase 3 would require full substantive compliance for all companies.³⁵

DPDP and Its Implications

The DPDP Act applied to any entity worldwide that processed the digital personal data of Indian residents in connection with offering goods or services in India. For global technology companies this meant that serving Indian customers at scale now triggered compliance obligations where the penalty for a single infraction could reach \$30 million.³⁶

The most operationally demanding provision for large technology companies was the Significant Data Fiduciary (SDF) designation, which the central government could apply to any data fiduciary based on the volume and sensitivity of data processed, the risk to data principals’ rights, and the potential impact on India’s “sovereignty, security, and integrity.”³⁷ Companies designated as SDFs were required to appoint an India-based Data Protection Officer, conduct annual Data Protection Impact Assessments (DPIAs) reviewed by independent auditors, report DPIA findings to the DPBI, and ensure that any algorithmic or technical systems they deployed were audited to prevent risks to data principal rights.³⁸ For AI companies deploying recommendation engines, predictive models, or training infrastructure that relied on large volumes of Indian user data, SDF designation was a near-certainty.³⁹

The most technically consequential and least resolved dimension of India’s regulatory framework was the DPDP Act’s right to erasure given the fundamental architecture of modern AI systems.⁴⁰ No major data protection regime had yet satisfactorily resolved this challenge, and India’s particular combination of large-scale AI ambitions and a consent-first legal framework made it especially acute. While the concept of “machine unlearning” was straightforward in a traditional database – locate rows and delete them – it was not in large language models (LLMs). LLMs did not store data in discrete

rows. Training an LLM involved ingesting a large corpus of data, masking elements of it so the model could try (and eventually learn) to predict the correct completion by repeating the process millions or billions of times and adjusting billions of probability weights. This process diffused information and data throughout the model's architecture.⁴¹

The effort to selectively remove or neutralize the influence of specific data points from already-trained AI models spawned two approaches. Exact unlearning required deleting the original data and fully retraining the model entirely from scratch, which fully removed data but could cost tens of millions of dollars, months of compute time, and generate substantial carbon emissions.⁴² Approximate unlearning adjusted a model's internal weights algorithmically to reduce the influence of specific data without full retraining.⁴³ One technology lawyer advising on the DPDP framework noted, "erasing influence from trained models is technically impossible, but compliance is not"—arguing that the practical standard should be stopping future use of the data, deleting identifiable inputs from training pipelines, and retraining from earlier checkpoints where feasible.⁴⁴

The Act reflected a philosophy that data was a national asset, and that sovereign control over how it flowed—both within India and across its borders—was integral to digital independence. The government determined which countries were safe recipients of Indian personal data, which entities became SDFs subject to heightened scrutiny, and whether specific processing activities were exempt on national security grounds.⁴⁵ At the same time, the government carved out broad exemptions for its own data processing, including activities related to national security, sovereignty, and public order.⁴⁶

Sovereign AI

India wasn't the only country to treat data, and now AI, as a strategic asset. Huang had gone so far as to assert that, "Every country needs to own the production of their intelligence."⁴⁷ Substantial AI sovereignty could require control across at least four distinct layers.

The Chip Layer and Hardware Supply Chain Layer

The foundation of an AI system was silicon, and the supply chain for advanced AI chips was geographically and commercially concentrated. Three US companies—NVIDIA, Broadcom, and AMD—accounted for the dominant share of advanced AI chip design.⁴⁸ TSMC in Taiwan was the dominant manufacturer of the sub-5nm process nodes required for frontier AI training and inference.⁴⁹ High-bandwidth memory (HBM)—the specialized memory architecture that enabled AI accelerators to process data at the speeds required for training—was produced by just three companies: Samsung and SK Hynix in South Korea, and Micron in the United States.⁵⁰ A single Nvidia Blackwell GPU chip involved design in California, fabrication at TSMC in Taiwan, HBM from South Korea, advanced packaging at TSMC's facilities in Taiwan.⁵¹ For any country seeking to develop AI, this supply chain created a fundamental exposure: access could be politically severed.

India faced this potentiality. IndiaAI Mission built a GPU fleet of over 38,000 chips by mid-2025—almost entirely composed of NVIDIA hardware procured through public-private partnerships with domestic industrial giants Larsen & Toubro (L&T), Yotta, and E2E Networks.⁵² This dependence reflected the current absence of any commercially viable alternative to NVIDIA's CUDA ecosystem for frontier AI training. Semiconductor Mission, announced as part of the government's broader industrial policy, targeted 20–25% participation in the global semiconductor value chain by 2047.⁵³

The Training Data Layer

Another dimension of AI sovereignty concerned the data on which models were trained. Knowledge, language, and cultural context were encoded in the system's learned weights. Frontier models trained primarily on English- or European-language internet content embedded implicit cultural and legal assumptions. The IndiaAI Mission's foundational models pillar—funding indigenous LLMs such as BharatGen and Sarvam-1, a two-billion-parameter model supporting ten major Indian languages—was motivated by this belief: sovereignty over AI's cultural and linguistic outputs required sovereignty over the training corpus.⁵⁴

The training data question extended to a second, commercially sensitive problem: the ongoing use of user interactions with deployed models as training signal for future model generations. A September 2025 Stanford study analyzed the privacy policies of six leading US AI foundation model developers—Amazon (Nova), Anthropic (Claude), Google (Gemini), Meta (Meta AI), Microsoft (Copilot), and OpenAI (ChatGPT)—and found that all six used consumer chat interactions to train their AI systems by default.⁵⁵ Enterprise and API customers were typically excluded from training data by default—creating a two-tier system in which free and low-cost consumer users subsidized model improvement with their data unless they opted out.

For India, this created a specific sovereignty concern. The hundreds of millions of Indian users engaging with free tiers of foreign LLMs were generating training data that flowed back to US AI companies and improved those companies' models. Each user who posed a question about Indian tax law, agricultural practices, or medical symptoms in an Indian language was providing labeled, domain-specific, culturally grounded data to a foreign model developer. That data—aggregated across hundreds of millions of interactions—could materially improve the foreign models' capabilities in domains where Indian developers were trying to establish a domestic advantage. Under the DPDP Act's framework, this practice was legally precarious.⁵⁶

The Model and Weights Layer

There was a meaningful sovereignty distinction between proprietary models accessible only via API—where the underlying weights, architecture, and fine-tuning data were entirely opaque and controlled by the developer—and open-weight models such as Meta's Llama series, which could be downloaded, self-hosted, and fine-tuned on domestic infrastructure. India's strategic posture had leaned toward open-weight architectures. A model whose weights were held domestically could not be switched off by a foreign company, and its behavior could be audited, adjusted, and aligned to Indian legal and cultural requirements.

The open-weight approach was not without its own limitations, however. The pre-training corpus, architecture decisions, and alignment training embedded in the base weights remained the product of foreign developers' choices. True model sovereignty required training frontier models from scratch on domestically curated, consented, and governed data—precisely the ambition behind BharatGen and the IndiaAI Datasets Platform, and the reason those initiatives were allocated nearly \$240 million (INR ₹2,000 crore) within the broader IndiaAI Mission budget.⁵⁷

The Data Center Layer and Access Continuity Layer

Many countries worried especially about where AI inference was taking place when a user prompted an AI model for a response. AI workloads processed in data centers owned or operated by foreign hyperscalers—for example, Amazon Web Services, Microsoft Azure, Google Cloud—were subject to the laws of those companies' home jurisdictions. A nation routing its government and other

AI workloads through US-headquartered cloud infrastructure had accepted, whether knowingly or not, that its data could be subject to American legal process.⁵⁸

IndiaAI Mission had directed compute investments toward domestically incorporated cloud providers operating India-domiciled data centers, rather than simply purchasing cloud credits from global hyperscalers. L&T's planned AI factory infrastructure—30 megawatts in Chennai and a new 40-megawatt facility in Mumbai—was positioned as sovereign cloud infrastructure capable of hosting government workloads outside foreign legal jurisdiction.⁵⁹ Still, the hardware inside those facilities was NVIDIA-designed, TSMC-fabricated, and HBM-sourced.

A further vulnerability was access continuity: the risk that a foreign-owned model or service could simply be turned off. Institutions that had integrated a proprietary foreign LLM into their operations — via API or hosted service — were dependent on the continued commercial availability, pricing stability, and geopolitical neutrality of that provider.

Firm Sovereignty

National leaders weren't the only ones worrying about control over their "intelligence." Corporate leaders also were paying much attention to where data came from for company AI uses and to where data went. They worried about explicit data that might be placed by employees via prompt into a chatbot, but also the data which could be abstracted from it. If many of their employees were asking about a certain question, would that give away their strategy to competitors? Would it make it easier for AI providers to become competitors themselves?

Firms faced contractual choices in addressing these concerns. Many companies and other organizations would sign enterprise agreements with large model providers. These, they hoped, would avoid the kinds of issues Samsung faced in 2023 for example, when employees had uploaded "sensitive code" to ChatGPT without confidentiality protections.⁶⁰ As the enterprise market developed, in exchange for their fees, firms could get guarantees about their prompts and their data. Anthropic promised, "By default, we will not use your inputs or outputs from our commercial products to train our models." Individual consumer (not commercial) users could modify their model improvement settings by toggling off the "Help Improve Claude" function in their privacy settings.⁶¹ These enterprise agreements would also solidify commitments about data retention and deletion. OpenAI offered, for example a "zero data retention" option where, "request and response bodies are not persisted to any logging mechanism."⁶² The model providers would also make "isolation guarantees", so that data would not be shared across customers.

Firms were not always satisfied to rely on these guarantees and made architectural choices to enhance their control. "Sixty-one percent of Western European CIOs and IT leaders said geopolitical factors will increase their reliance on local or regional cloud providers," according to a November 2025, Gartner, Inc. survey.⁶³ The company predicted "more than 75% of all enterprises outside of the US [would] have a digital sovereignty strategy" by 2030. Some companies went further and chose to deploy open-weight models on private infrastructure.

Firms worried about what might be considered an AI version of "Sherlocking." The term had caught on in the 1990s, when Apple was said to have launched features in its search app named Sherlock similar to those provided by a competing third-party product in its app store, Watson.⁶⁴ In the AI context, "Sherlocking" might apply to giant AI companies launching products or features that other companies had built, or even contemplated building, on top of their tools. Startup founders were warned that maybe the giant AI companies were "creating playgrounds for other startups, which they

then mine for data, ideas and inspiration.”⁶⁵ Even if model providers were not training their models on prompts or data, the belief was that they could still divine much from the overall shape of usage.

Firms thus appeared to face related concerns as nations. Microsoft’s CEO Satya Nadella had explicitly extended this idea of national sovereignty to firms in comments at the World Economic Forum in Davos in January 2026. He said “What does sovereignty mean? In the AI world, the topic that is least talked about, but I feel will be most talked about in this calendar year will be the sovereignty of a firm. Just imagine if [at] your firm, you’re not able to embed the tacit knowledge of the firm in a set of weights in a model you control. By definition you have no sovereignty. That means you are leaking enterprise value to some model company somewhere...Control of destiny means that you have the ability to produce something that is unique that is preserved...That is what’ll give you real sovereignty.”⁶⁶

Sovereignty in the Age of Agents

As reliance on AI evolved from models to agents, views about AI sovereignty did, too. The first wave of AI sovereignty concerns were premised primarily on a world in which AI systems responded to queries. In early 2026, the industry appeared to be moving materially towards agentic AI: autonomous systems that did not merely generate responses but perceived their environment, decomposed goals into multi-step plans, invoked external tools, wrote and executed code, accessed files and APIs, sent communications, and completed tasks with minimal human steering.⁶⁷ This shift—from AI as a tool humans operated to AI as agents that operated on a humans’ behalf—would alter sovereignty concerns from questions about where data was stored and perhaps “seen” by a few major AI companies to questions about what an autonomous system was doing with that data at every moment of its operation across every system it touches.

Structural Data Sovereignty in Agentic Systems

Agentic AI systems, by design, moved data between systems in ways that could be difficult to observe, log, attribute, and govern. Where a traditional application made a defined API call with a predictable data payload, an AI agent made dynamic, context-dependent, probabilistic decisions about what data to retrieve, where to send it, what to include in a prompt to an external model, and what actions to take with the results. Researchers had documented what security expert Simon Willison termed the “lethal trifecta” of AI agent risk: access to private data, exposure to untrusted content, and the ability to communicate externally.⁶⁸ When these three capabilities combined—as they did in many agentic deployments—attackers could trick an agent into accessing private information and exfiltrating it, often without any alert being triggered.

Analysts found that AI agents moved 16 times more data than human users, running continuously and chaining tasks across multiple SaaS applications with no natural stopping points.⁶⁹ A single improperly configured Glean agent was documented downloading over 16 million files.⁷⁰ *Cisco’s State of AI Security 2026* report found that only 29% of organizations deploying agentic AI reported readiness to secure those deployments, and that 80% had already encountered risky agent behaviors including unauthorized system access and improper data exposure.⁷¹ A McKinsey analysis of agentic AI risks called out “untraceable data leakage” as a category in which “autonomous agents exchanging data without oversight obscure leaks and evade audits.”⁷²

For India, agentic systems created a new and different sovereignty exposure. A foreign-hosted LLM that received a query and returned a response was a relatively bounded interaction—the data that flowed to the provider was the text of the query and perhaps any data sent with it. An agentic system

authorized to act on behalf of a government official, a corporate executive, or a healthcare professional could autonomously read documents, query databases, compose and send communications, make bookings, execute transactions, and chain all of these actions across multiple foreign-owned cloud services before the human principal was even aware of what occurred. Much of this would seem subject to DPDP Act obligations, and much of it at the time was processed through foreign infrastructure.

Multi-agent systems—in which specialized agents orchestrated and communicated with each other to complete complex tasks—amplified these risks further. A compromised research agent could insert hidden instructions into output consumed by a financial agent, which then executed unintended transactions.⁷³ A scheduling agent in a healthcare system could falsely escalate a request as originating from a licensed physician, tricking a clinical-data agent into releasing patient records without triggering security alerts.⁷⁴ The OWASP GenAI Security Project, which released its Top 10 Risks for Agentic Applications in December 2025, cataloged these failure modes—“prompt injection, tool misuse, privilege escalation, memory poisoning, and cascading failures across agent networks”—as the defining security challenges of the agentic era.⁷⁵

Claws and the Speed of Agentic Adoption

One catalyst for the industry’s reckoning with agents was the take-up of “claws” in early 2026. Claws were described as, “a new layer on top of LLM agents.”⁷⁶ Claws could “plan and execute tasks on their own and...spin up their own subagents to tackle specialized tasks, access files and themselves delegate tasks to other subagents,” one March *Wall Street Journal* article explained.⁷⁷ It added, “they represent a big leap beyond question-and-answer style AI chatbots as well as recent iterations of AI agents, which typically have narrow use cases and run for a set amount of time.”⁷⁸ OpenClaw, the most famous among them, would by default, “wake itself up” every 30 minutes and “ask itself: Is there anything I should be doing right now?”⁷⁹ It could “check reminders, follow up on loose threads, and review inboxes” in ways that made it feel “proactive instead of reactive.”⁸⁰

OpenClaw, created by Peter Steinberger and launched in December 2025 under the name Clawdbot (later rebranded as Moltbot and then OpenClaw following trademark disputes), was open-source and self-hosted.⁸¹ Its code was entirely available for viewing and modification via GitHub.⁸² OpenClaw could be hosted on personal computers. Its popularity had led to a run on Mac Minis, in part because the computers could be packed with enough memory to run OpenClaw’s agents effectively, and left powered-on all the time so they could run persistently.⁸³ The project became the fastest-growing open-source repository in GitHub history.⁸⁴ (See **Exhibit 6** for OpenClaw stars history). Huang called it “probably the single most important release of software, you know, probably ever.”⁸⁵ In February, OpenAI announced Steinberger would join the company, that the open-source tool would “live in a foundation” inside of it, and that the tool would “quickly become core to our product offerings.”⁸⁶

Among OpenClaw’s promised appeal, elements of AI sovereignty were a key feature. Its developers promoted the fact that it “runs on your machine.” “Mac, Windows, or Linux. Anthropic, OpenAI or local models,” they wrote.⁸⁷ “Private by default – your data stays yours.” Users could configure OpenClaw to run with AI models of their own choosing and even, if they wanted, with models small enough to run on their own computer. In China, leaders were pushing adoption of OpenClaw in part, it was said, so they could bet “big on lower-cost domestic AI models” and create a “parallel ecosystem” to challenge “Western dominance in the next wave of artificial intelligence.”⁸⁸

In India and elsewhere, an OpenClaw-driven “DIY^a AI Agent Boom” was also said to be underway, in part because of the “appeal” of its “autonomy”, but against the backdrop of significant security risks.⁸⁹ It connected to messaging platforms including WhatsApp, Telegram, Slack, and Discord; executed shell commands; could read and write files, browse the web, manage calendars and email, send messages autonomously; and supported community-built “skills” distributed through its marketplace ClawHub.⁹⁰ When run locally, an OpenClaw agent with typical integrations had persistent access to a user’s digital life and, by extension if deployed in an enterprise context, to the organization’s. Researchers at Cisco described OpenClaw as “groundbreaking” from a capability perspective but “an absolute nightmare” from a security perspective.⁹¹ Separately, researchers found that the Moltbook platform (a social network built around OpenClaw agents) had an unsecured database exposing 35,000 email addresses and 1.5 million agent API tokens; and that approximately 12% of all skills in ClawHub were malicious, many of them serving as delivery vectors for malware.⁹² Microsoft published security guidance recommending that any evaluation occur only in a fully isolated environment with dedicated, non-privileged credentials and no access to sensitive data.⁹³ Infamously, a Meta safety researcher experimenting with OpenClaw posted about how the tool planned to erase all of her emails, was refusing to stop and that she had to, “RUN to my Mac mini like I was defusing a bomb.”⁹⁴

NemoClaw: NVIDIA’s Response?

Ahead of GTC, word began to spread that NVIDIA was readying to embrace claws well beyond potential internal usage at the company. *WIRED* reported that NVIDIA had begun pitching a new offering, “NemoClaw,” to major software companies. “The platform will allow these companies to dispatch AI agents to perform tasks for their own workforces,” two reporters foreshadowed.⁹⁵ “Companies will be able to access the platform regardless of whether their products run on Nvidia’s chips,” they’d been told.⁹⁶ NVIDIA was said to be planning to “offer security and privacy tools as part of this new open-source agent platform.” The early thinking was that NVIDIA was aiming to “court enterprise software companies by offering additional layers of security for AI agents.”⁹⁷ It would also be “another step in the company’s embrace of open-source models.”⁹⁸ NVIDIA had already launched their own open models, dubbed Nemotron, and in December 2025 had noted, “NVIDIA Nemotron supports broader sovereign AI efforts, with organizations from Europe to South Korea adopting open, transparent, efficient models that allow them to build AI systems aligned to their own data, regulations, and values.”⁹⁹ The company had said that early corporate Nemotron adopters included: Accenture, CrowdStrike, Cursor, Deloitte, EY, Palantir, Perplexity, Zoom, and others.¹⁰⁰

Now, companies – and countries – would be watching for more details of NemoClaw and wondering to what extent it would offer a meaningful answer to AI sovereignty challenges. India’s DPDP Act was not drafted with agentic systems in mind. The IndiaAI Governance Guidelines did not address agents and certainly not claws. This regulatory gap was not unique to India. Most AI governance frameworks did not “address the specific technical controls needed for agentic deployments, such as tool call parameter validation, prompt injection logging, or containment testing for multi-agent systems.”¹⁰¹ But the gap carried perhaps particular urgency in India because the country’s concurrent ambitions—to deploy AI at scale, to train sovereign foundational models on Indian data, and to protect Indian citizens’ personal data under the DPDP Act—all now potentially rested on agentic systems as the primary execution layer. Claws, not chatbots, might be the interface

^a “Do it Yourself”

through which hundreds of millions of Indians interacted with AI-powered government services, financial systems, healthcare platforms, and more.

One could assume the latest news on NemoClaw had been carefully timed on the heels of the OpenClaw craze and just ahead of when Huang would take center stage at GTC in a week. NemoClaw would be “part of a broader strategy to maintain its dominance in AI infrastructure at a time when leading AI labs are building their own custom chips,” the *WIRED* article noted, suggesting the stakes.¹⁰² “Nvidia’s software strategy until now had been heavily reliant on its CUDA platform, a famously proprietary system that locks developers into building software for Nvidia’s GPUs and has created a crucial ‘moat’ for the company.”¹⁰³ How, architecturally, would NemoClaw address the security and sovereignty concerns in a new age of AI agents and claws? And how big a bet would Huang place on it at GTC? More and more it seemed that countries and companies wanted to bring the AI to their data instead of bringing their data to AI. How far would Huang help them in each direction?

Exhibit 1 NVIDIA Historical Stock Price - April 21, 2014 to February 28, 2026

Source: NVIDIA. "Stock Chart." <https://investor.nvidia.com/stock-info/stock-quote-and-chart/default.aspx>, accessed March 2026.

Exhibit 2 NVIDIA Income Statement

	Year Ended		
	Jan 25, 2026	Jan 26, 2025	Jan 28, 2024
Revenue	\$ 215,938	\$ 130,497	\$ 60,922
Cost of revenue	62,475	32,639	16,621
Gross profit	153,463	97,858	44,301
Operating expenses			
Research and development	18,497	12,914	8,675
Sales, general and administrative	4,579	3,491	2,654
Total operating expenses	23,076	16,405	11,329
Operating income	130,387	81,453	32,972
Interest income	2,300	1,786	866
Interest expense	(259)	(247)	(257)
Other income, net	9,022	1,034	237
Total other income, net	11,063	2,573	846
Income before income tax	141,450	84,026	33,818
Income tax expense	21,383	11,146	4,058
Net income	\$ 120,067	\$ 72,880	\$ 29,760
Net income per share:			
Basic	\$ 4.93	\$ 2.97	\$ 1.21
Diluted	\$ 4.90	\$ 2.94	\$ 1.19

Source: Casewriters adapted from NVIDIA 10-K for the year ended January 2026. <https://investor.nvidia.com/financial-info/sec-filings/sec-filings-details/default.aspx?FilingId=19184805>, accessed March 2026.

Exhibit 3a NVIDIA Architectures


Blackwell (March 2024)
Hopper (March 2022)
Ada Lovelace (September 2022)

Previous Architectures:

Ampere (2020)
Turing (2018)
Volta (2017)
Pascal (2016)
Maxwell (2014)
Kepler (2012)
Fermi (2010)
Tesla (2006)
Curie (2004)
Rankine (2003)
Kelvin (2001)
Celsius (1999)

Source: Casewriters adaptation from NVIDIA, "Technologies" <https://www.nvidia.com/en-us/technologies/>, accessed March 2026.

Exhibit 3b NVIDIA Blackwell Products

**NVIDIA GB300 NVL72**


The NVIDIA GB300 NVL72 delivers unparalleled AI reasoning inference performance, featuring 65X more AI compute than Hopper systems.

[Learn More >](#)

**NVIDIA DGX SuperPOD**

NVIDIA DGX SuperPOD™ is a turnkey AI data center solution that delivers leadership-class accelerated infrastructure with scalable performance for the most demanding AI training and inference workloads.

[Learn More >](#)

**NVIDIA RTX PRO in the Data Center**


Deliver powerful AI and graphics acceleration, essential enterprise features, and the flexibility to handle a wide range of workloads, from agentic and physical AI to visual computing and virtual workstations accelerated by NVIDIA RTX PRO™ data center GPUs.

[Learn More >](#)

**NVIDIA RTX PRO Workstations**

Bring the latest breakthroughs in AI, ray tracing, and neural graphics technology to power the most innovative workflows in design, engineering, and beyond with NVIDIA RTX PRO GPUs.

[Learn More >](#)

**NVIDIA DGX Station**

Unlike any AI desktop computer before, this system features NVIDIA Blackwell GPUs, the Grace CPU Superchip, and large coherent memory, delivering unparalleled compute performance.

[Learn More >](#)

**NVIDIA DGX Spark**

A compact, personal AI supercomputer with the NVIDIA GB10 Grace Blackwell Superchip, delivering high-performance AI capabilities and support for models up to 200 billion parameters.

[Learn More >](#)

**NVIDIA HGX B300**


NVIDIA HGX™ B300 is built for the age of AI reasoning with enhanced compute and increased memory.

[Learn More >](#)

**NVIDIA GB200 NVL72**

The NVIDIA GB200 NVL72 connects 36 NVIDIA Grace CPUs and 72 NVIDIA Blackwell GPUs in a rack-scale, liquid-cooled design.


[Learn More >](#)

**NVIDIA GB200 NVL4**

Purpose-built for scientific computing, the NVIDIA GB200 NVL4 unlocks the future of converged high-performance computing and AI.

[Learn More >](#)

Source: NVIDIA. "NVIDIA Blackwell Products", <https://www.nvidia.com/en-us/data-center/technologies/blackwell-architecture/?ncid=no-ncid>, accessed March 2026.

Exhibit 4 NVIDIA's CUDA


What Is CUDA?

CUDA is NVIDIA's platform for accelerated computing, providing the software layer that enables applications to harness the power of GPUs. Developers can program in languages such as C++, Python, and Fortran or leverage GPU-accelerated libraries and frameworks like PyTorch. This flexibility lets developers integrate GPU computing into any layer of their software stack to achieve optimal functionality and performance.

The CUDA Toolkit, an integral component of the CUDA platform, provides the compiler, libraries, and developer tools required to develop GPU applications.


Source: NVIDIA. "What is CUDA", <https://developer.nvidia.com/cuda>, accessed March 2026.

Exhibit 5 M.A.N.A.V.



Source: "PM's address at India AI Impact Summit 2026." Prime Minister's Office. https://www.pmindia.gov.in/en/news_updates/pms-address-at-india-ai-impact-summit-2026/, accessed March 2026

Exhibit 6 OpenClaw History



Source: Star History. <https://www.star-history.com/?repos=openclaw%2Fopenclaw&type=date&legend=top-left>, accessed March 2026.

Endnotes

- ¹ NVIDIA, “Efficiency isn’t just about better tools anymore – it’s about autonomous teammates.” X.com, March 13, 2026. <https://x.com/nvidia/status/2032592951424012313>, accessed March 2026.
- ² Prime Minister’s Office, “PM’s Address at India AI Impact Summit 2026.” Prime Minister’s Office, February 19, 2026. https://www.pmindia.gov.in/en/news_updates/pms-address-at-india-ai-impact-summit-2026/, accessed March 2026.
- ³ Niket Nishant and Rashika Singh, “Nvidia hits \$5 trillion valuation as AI boom powers meteoric rise.” *Reuters*, October 29, 2025. <https://www.reuters.com/business/nvidia-poised-record-5-trillion-market-valuation-2025-10-29/>, accessed March 2026.
- ⁴ NVIDIA, “NVIDIA Announces Financial Results for Fourth Quarter and Fiscal 2026.” *NVIDIA Press Release*, February 25, 2026. <https://investor.nvidia.com/news/press-release-details/2026/NVIDIA-Announces-Financial-Results-for-Fourth-Quarter-and-Fiscal-2026/default.aspx>, accessed March 2026.
- ⁵ NVIDIA, “NVIDIA Announces Financial Results for Fourth Quarter and Fiscal 2026.” *NVIDIA Press Release*, February 25, 2026. <https://investor.nvidia.com/news/press-release-details/2026/NVIDIA-Announces-Financial-Results-for-Fourth-Quarter-and-Fiscal-2026/default.aspx>, accessed March 2026.
- ⁶ NVIDIA, “A Timeline of Innovation.” *NVIDIA, About Us*. <https://www.nvidia.com/en-us/about-nvidia/corporate-timeline/>, accessed March 2026.
- ⁷ Sequoia Capital, “Crucible Moments: Nvidia – Jensen Huang.” *Sequoia Capital podcast*. <https://www.youtube.com/watch?v=75RyknkIKkw>, accessed March 2026.
- ⁸ Sequoia Capital, “Crucible Moments: Nvidia – Jensen Huang.” *Sequoia Capital podcast*. <https://www.youtube.com/watch?v=75RyknkIKkw>, accessed March 2026.
- ⁹ Sequoia Capital, “Crucible Moments: Nvidia – Jensen Huang.” *Sequoia Capital podcast*. <https://www.youtube.com/watch?v=75RyknkIKkw>, accessed March 2026.
- ¹⁰ Vivek Singh, “Game-Change: How the World’s First GPU Leveled Up Gaming and Ignited the AI Era.” *NVIDIA Company Blog*, October 11, 2025. <https://blogs.nvidia.com/blog/first-gpu-gaming-ai/>, accessed March 2026.
- ¹¹ NVIDIA, “2022 Annual Report.”, *NVIDIA Annual Report*, 2022. https://s201.q4cdn.com/141608511/files/doc_financials/2022/ar/2022-annual-report-insert-web-IR-update.pdf, accessed March 2026. p 148.
- ¹² Vivek Singh, “Game-Change: How the World’s First GPU Leveled Up Gaming and Ignited the AI Era.” *NVIDIA Company Blog*, October 11, 2025. <https://blogs.nvidia.com/blog/first-gpu-gaming-ai/>, accessed March 2026.
- ¹³ NVIDIA, “A Timeline of Innovation.” *NVIDIA, About Us*. <https://www.nvidia.com/en-us/about-nvidia/corporate-timeline/>, accessed March 2026.
- ¹⁴ Steven Rosenbush and Isabelle Bousquette, “Nvidia CEO Says AI Computing Needs to Surge 100 Fold.” *The Wall Street Journal*, March 18, 2025. <https://www.wsj.com/tech/ai/nvidia-ceo-says-ai-computing-needs-to-surge-100-fold-at-ai-super-bowl-event>, accessed March 2026.
- ¹⁵ Blake Crosley, “Inference-Time Scaling: The New Training Frontier for AI Reasoning.” *Introl blog*, December 12, 2025. <https://introl.com/blog/inference-time-scaling-research-reasoning-models-december-2025>, accessed March 2026.
- ¹⁶ Cade Metz and Karen Weise, “What Exactly are A.I. Companies Trying to Build? Here’s a Guide.” *The New York Times*. <https://www.nytimes.com/2025/09/16/technology/what-exactly-are-ai-companies-trying-to-build-heres-a-guide.html>, accessed March 2026.
- ¹⁷ xAI. “Colossus: Our Gigafactory of Compute.” *xAI website*. <https://x.ai/colossus>, accessed March 2026.
- ¹⁸ Brendan Bordelon and Gabby Miller. “The big split driving the tricky politics of AI data centers.” *POLITICO*, February 6, 2026. <https://www.politico.com/news/2026/02/06/tech-industry-ai-data-centers-politics-00762348>, accessed March 2026.
- ¹⁹ Cade Metz and Karen Weise, “What Exactly are A.I. Companies Trying to Build? Here’s a Guide.” *The New York Times*. <https://www.nytimes.com/2025/09/16/technology/what-exactly-are-ai-companies-trying-to-build-heres-a-guide.html>, accessed March 2026.

- ²⁰ Harsh Chauhan, "Nvidia's AI Dominance: Data Center Revenue Poised for 165% Surge by 2027." *The Motley Fool*, November 8, 2025. <https://finance.yahoo.com/news/nvidias-ai-dominance-data-center>, accessed March 2026.
- ²¹ NVIDIA. "NVIDIA Blackwell Platform Arrives to Power a New Era of Computing." *NVIDIA Newsroom*. March 18, 2024. <https://nvidianews.nvidia.com/news/nvidia-blackwell-platform-arrives-to-power-a-new-era-of-computing>, accessed March 2026.
- ²² NVIDIA. "NVIDIA Blackwell Platform Arrives to Power a New Era of Computing." *NVIDIA Newsroom*. March 18, 2024. <https://nvidianews.nvidia.com/news/nvidia-blackwell-platform-arrives-to-power-a-new-era-of-computing>, accessed March 2026.
- ²³ Laura Bratton, "Nvidia's Big Tech customers might also be its biggest competitive threat." *Yahoo! Finance*, October 20, 2025. <https://finance.yahoo.com/news/nvidias-big-tech-customers-might-also-be-its-biggest-competitive-threat-153032596.html>, accessed March 2026.
- ²⁴ NVIDIA, "2026 10-K." *NVIDIA Form 10-K*, 2026. <https://d18rn0p25nwr6d.cloudfront.net/CIK-0001045810/e361e58a-7483-44f5-bc62-a9080ae6ec72.pdf>, accessed March 26. P23.
- ²⁵ NVIDIA, "2026 10-K." *NVIDIA Form 10-K*, 2026. <https://d18rn0p25nwr6d.cloudfront.net/CIK-0001045810/e361e58a-7483-44f5-bc62-a9080ae6ec72.pdf>, accessed March 26. P40.
- ²⁶ Melissa Otto, "Nvidia post earnings review: Q4 2026 and outlook." *S&P Global Market Intelligence*, March 12, 2026. <https://www.spglobal.com/market-intelligence/en/news-insights/research/2026/03/nvidia-post-earnings-review-q4-2026-and-outlook>, accessed March 2026.
- ²⁷ Prime Minister's Office, "PM's Address at India AI Impact Summit 2026." Prime Minister's Office, February 19, 2026. https://www.pmindia.gov.in/en/news_updates/pms-address-at-india-ai-impact-summit-2026/, accessed March 2026.
- ²⁸ Prime Minister's Office, "PM's Address at India AI Impact Summit 2026." Prime Minister's Office, February 19, 2026. https://www.pmindia.gov.in/en/news_updates/pms-address-at-india-ai-impact-summit-2026/, accessed March 2026.
- ²⁹ Prime Minister's Office, "PM's Address at India AI Impact Summit 2026." Prime Minister's Office, February 19, 2026. https://www.pmindia.gov.in/en/news_updates/pms-address-at-india-ai-impact-summit-2026/, accessed March 2026.
- ²⁹ Anirudh Burman, "Understanding India's New Data Protection Law." *Carnegie India*, October 3, 2023. <https://carnegieendowment.org/research/2023/10/understanding-indias-new-data-protection-law>, accessed March 2026.
- ³⁰ "The Information Technology (Reasonable Security Practices and Procedures and Sensitive Personal Data or Information) Rules, 2011, India. *WIPO Lex*, April 11, 2011. <https://www.wipo.int/wipolex/en/legislation/details/15063>, accessed March 2016.
- ³¹ Jyoti Panday, "India's Supreme Court Upholds Right to Privacy as a Fundamental Right – and It's About Time. *Electronic Frontier Foundation*, August 28, 2017. <https://www.eff.org/deeplinks/2017/08/indias-supreme-court-upholds-right-privacy-fundamental-right-and-its-about-time>, accessed March 2026.
- ³² Anirudh Burman, "Understanding India's New Data Protection Law." *Carnegie India*, October 3, 2023. <https://carnegieendowment.org/research/2023/10/understanding-indias-new-data-protection-law>, accessed March 2026.
- ³³ Anirudh Burman, "Understanding India's New Data Protection Law." *Carnegie India*, October 3, 2023. <https://carnegieendowment.org/research/2023/10/understanding-indias-new-data-protection-law>, accessed March 2026.
- ³⁴ DLA Piper, "Data Protection Laws in India." *DLA Piper*. <https://www.dlapiperdataprotection.com/?t=law&c=IN>, accessed March 2026.
- ³⁵ DLA Piper, "Data Protection Laws in India." *DLA Piper*. <https://www.dlapiperdataprotection.com/?t=law&c=IN>, accessed March 2026.
- ³⁶ ITIF, "India's Personal Data Protection Regulation." *Information Technology & Innovation Foundation*, February 5, 2026. <https://itif.org/publications/2025/05/14/india-personal-data-protection-regulation/>, accessed March 2026.
- ³⁷ Latham & Watkins, "India's Digital Personal Data Protection Act 2023 vs. the GDPR: A Comparison." *Latham & Watkins*, December 2023. <https://www.lw.com/admin/upload/SiteAttachments/Indias-Digital-Personal-Data-Protection-Act-2023-vs-the-GDPR-A-Comparison.pdf>, accessed March 2026.
- ³⁸ ITIF, "India's Personal Data Protection Regulation." *Information Technology & Innovation Foundation*, February 5, 2026. <https://itif.org/publications/2025/05/14/india-personal-data-protection-regulation/>, accessed March 2026.

- ³⁹ Aakanksha Tewari, “DPDPA’s Global Reach: Cross-Border Data, AI Impact, and International Alignment.” *TrustArc*. <https://trustarc.com/resource/dpdpa-global-reach-cross-border-data-ai/>, accessed March 2026.
- ⁴⁰ Indian Ministry of Law and Justice. “The Digital Personal Data Protection Act, 2023.” August 11, 2023. <https://www.meity.gov.in/static/uploads/2024/06/2bf1f0e9f04e6fb4f8fef35e82c42aa5.pdf>, accessed March 2026.
- ⁴¹ Vikrant Rana, Anuradha Gandhi, and Prateek Chandgothia. “Effect of Digital Personal Data Protection Rules, 2025 on AI Regulation.” S.S. Rana & Co, November 24, 2025. <https://ssrana.in/articles/effect-of-digital-personal-data-protection-rules-2025-on-ai-regulation/>, accessed March 2026.
- ⁴² Tanusha Tyagi, “Rethinking the Right to Be Forgotten in India’s AI Governance.” *Observer Research Foundation*, January 12, 2026. <https://www.orfonline.org/expert-speak/rethinking-the-right-to-be-forgotten-in-india-s-ai-governance>, accessed March 2026.
- ⁴³ Tanusha Tyagi, “Rethinking the Right to Be Forgotten in India’s AI Governance.” *Observer Research Foundation*, January 12, 2026. <https://www.orfonline.org/expert-speak/rethinking-the-right-to-be-forgotten-in-india-s-ai-governance>, accessed March 2026.
- ⁴⁴ Prabhanu Kumas Das, “DPDP Rules And AI: How India’s Personal Data Framework Will Affect Data Collection And Retention.” *Medianama*, November 19, 2025. <https://www.medianama.com/2025/11/223-dpdp-rules-ai-indias-personal-data-framework-data-collection-retention/>, accessed March 2026.
- ⁴⁵ Indian Ministry of Law and Justice. “The Digital Personal Data Protection Act, 2023.” August 11, 2023. <https://www.meity.gov.in/static/uploads/2024/06/2bf1f0e9f04e6fb4f8fef35e82c42aa5.pdf>, accessed March 2026.
- ⁴⁶ Indian Ministry of Law and Justice. “The Digital Personal Data Protection Act, 2023.” August 11, 2023. <https://www.meity.gov.in/static/uploads/2024/06/2bf1f0e9f04e6fb4f8fef35e82c42aa5.pdf>, accessed March 2026.
- ⁴⁷ NVIDIA, “NVIDIA CEO: Every Country Needs Sovereign AI.” <https://blogs.nvidia.com/blog/world-governments-summit/>, accessed January 2026.
- ⁴⁸ TrendForce, “Global Top 10IC Design Houses See 49% YoY Growth in 2024.” *TrendForce*, March 17, 2025. <https://www.trendforce.com/presscenter/news/20250317-12519.html>, accessed March 2026.
- ⁴⁹ TrendForce, “TSMC’s Advanced Processes Remain Resilient Amid Challenges.” *TrendForce*, April 8, 2024, <https://www.trendforce.com/news/2024/04/08/news-tsmcs-advanced-processes-remain-resilient-amid-challenges/>, accessed March 2026.
- ⁵⁰ American Compass, “Stop Selling the Rope.” *American Compass*, October 27, 2025. <https://americancompass.org/stop-selling-the-rope/>, accessed March 2026.
- ⁵¹ Sourceability, “AI chip shortages deepen amid tariff risks.” *Sourceability*, June 2, 2025. <https://americancompass.org/stop-selling-the-rope/>, accessed March 2026.
- ⁵² Deepankshi Agnihotry, “IndiaAI Mission (2024)”, *Impact and Policy Research Institute*, March 6, 2025. <https://www.impriindia.com/insights/indiaai-mission-2024/>, accessed March 2026.
- ⁵³ AI & CloudSummit Team, “China Bans Nvidia AI Chips in Technology Sovereignty Showdown.” *AI & Cloud Summit*, September 18, 2025. <https://cloudsummit.eu/blog/china-bans-nvidia-ai-chips-technology-sovereignty-showdown>, accessed March 2026.
- ⁵⁴ PIB Headquarters, “Transforming India with AI.” *PIB Headquarters*, December 30, 2025. <https://www.pib.gov.in/PressReleasePage.aspx?PRID=2209737®=3&lang=2>, accessed March 2026.
- ⁵⁵ Jennifer King, Kevin Klyman, Emily Capstick, Tiffany Saade, and Victoria Hsieh, “User Privacy and Large Language Models: An Analysis of Frontier Developers’ Privacy Policies.” *arXiv*, September 5, 2025. <https://arxiv.org/html/2509.05382v1>, accessed March 2026.
- ⁵⁶ Bilal Mohamed, “Five Ways in Which the DPDPA Could Shape the Development of AI in India.” *Future of Privacy Forum*, September 6, 2024. <https://fpf.org/blog/five-ways-in-which-the-dpdpa-could-shape-the-development-of-ai-in-india/>, accessed March 2026.
- ⁵⁷ PIB Headquarters, “Transforming India with AI.” *PIB Headquarters*, December 30, 2025. <https://www.pib.gov.in/PressReleasePage.aspx?PRID=2209737®=3&lang=2>, accessed March 2026.

- ⁵⁸ Irena Mroz, "Understanding the U.S. CLOUD Act: Impact on Compliance, Agreement, and Data Protection." *archTIS*, May 28, 2025. <https://www.archtis.com/understanding-the-us-cloud-act/>, accessed March 2026.
- ⁵⁹ Jay Puri, "India Fuels Its AI Mission with NVIDIA." *NVIDIA Company Blog*, February 17, 2026. <https://blogs.nvidia.com/blog/india-ai-mission-infrastructure-models/>, accessed March 2026.
- ⁶⁰ Mark Gurman, "Samsung Bans Staff's AI Use After Spotting ChatGPT Data Leak." *Bloomberg*, May 1, 2023. <https://www.bloomberg.com/news/articles/2023-05-02/samsung-bans-chatgpt-and-other-generative-ai-use-by-staff-after-leak?embedded-checkout=true>, accessed March 2026.
- ⁶¹ Anthropic, "How do I change my model improvement privacy settings." Anthropic. <https://privacy.claude.com/en/articles/12109829-how-do-i-change-my-model-improvement-privacy-settings>, accessed March 2026.
- ⁶² Steve Charkoudian and Omer Tene, "Contracting around AI: Reading the fine print." *iapp*, November 27, 2024. <https://iapp.org/news/a/contracting-around-ai-reading-the-fine-print>, accessed March 2024.
- ⁶³ Gartner. "Gartner Survey Reveals Geopolitics Will Drive 61% of CIOs and IT Leaders in Western Europe to Increase Reliance on Local Cloud Providers." Gartner, Inc. November 12, 2025. <https://www.gartner.com/en/newsroom/press-releases/2025-11-12-gartner-survey-reveals-geopolitics-will-drive-61-percent-of-cios-and-information-technology-leaders-in-western-europe-to-increase-reliance-on-local-cloud-providers>, accessed March 2026.
- ⁶⁴ Ivan Mehta, "The apps that Apple sherlocked at WWDC 2024." *TechCrunch*, June 12, 2024. <https://techcrunch.com/2024/06/12/the-apps-that-apple-sherlocked-at-wwdc/>, accessed March 2026.
- ⁶⁵ Jason Calacanis, "Hear my WARNING for ANY APP using OpenAI's API." *LinkedIn*, November 2025. https://www.linkedin.com/posts/jasoncalacanis_hear-my-warning-for-any-app-using-openai-activity-7390797620165701632-ButX/, accessed March 2026.
- ⁶⁶ Satya Nadella and Laurence Fink, "Conversation with Satya Nadella, CEO of Microsoft." *World Economic Forum*, January 20, 2026. <https://www.weforum.org/meetings/world-economic-forum-annual-meeting-2026/sessions/conversation-with-satya-nadella-ceo-of-microsoft/>, accessed March 2026.
- ⁶⁷ Jeff Burt, "Nvidia Says OpenClaw is to Agentic AI What GPT was to Chattybots." *The NextPlatform*, March 17, 2026. <https://www.nextplatform.com/ai/2026/03/17/nvidia-says-openclaw-is-to-agentic-ai-what-gpt-was-to-chattybots/5209428>, accessed March 2026.
- ⁶⁸ Louis Columbus, "OpenClaw proves agentic AI works. It also proves your security model doesn't" *Venture Beat*, January 30, 2026. <https://venturebeat.com/security/openclaw-agentic-ai-security-risk-ciso-guide>, accessed March 2026.
- ⁶⁹ Obsidian Security Team, "The 2025 AI Agent Security Landscape: Players, Trends, and Risks." *Obsidian*, October 23, 2025. <https://www.obsidiansecurity.com/blog/ai-agent-market-landscape>, accessed March 2026.
- ⁷⁰ Obsidian Security Team, "The 2025 AI Agent Security Landscape: Players, Trends, and Risks." *Obsidian*, October 23, 2025. <https://www.obsidiansecurity.com/blog/ai-agent-market-landscape>, accessed March 2026.
- ⁷¹ Mirko Zorz, "Enterprises are racing to secure agentic AI deployments." *Help Net Security*, February 23, 2026. <https://www.helpnetsecurity.com/2026/02/23/ai-agent-security-risks-enterprise/>, accessed March 2026.
- ⁷² Benjamin Klein, Charlie Lewis, Rich Isenberg, Dante Gabrielli, Helen Möllering, Raphael Engler, and Vincent Yuan. "Deploying agentic AI with safety and security: A playbook for technology leaders." *McKinsey Quarterly*, October 16, 2025. <https://www.mckinsey.com/capabilities/risk-and-resilience/our-insights/deploying-agentic-ai-with-safety-and-security-a-playbook-for-technology-leaders>, accessed March 2026.
- ⁷³ Mirko Zorz, "Enterprises are racing to secure agentic AI deployments." *Help Net Security*, February 23, 2026. <https://www.helpnetsecurity.com/2026/02/23/ai-agent-security-risks-enterprise/>, accessed March 2026.
- ⁷⁴ Benjamin Klein, Charlie Lewis, Rich Isenberg, Dante Gabrielli, Helen Möllering, Raphael Engler, and Vincent Yuan. "Deploying agentic AI with safety and security: A playbook for technology leaders." *McKinsey Quarterly*, October 16, 2025. <https://www.mckinsey.com/capabilities/risk-and-resilience/our-insights/deploying-agentic-ai-with-safety-and-security-a-playbook-for-technology-leaders>, accessed March 2026.
- ⁷⁵ Scott Clinton, "OWASP GenAI Security Project Releases Top 10 Risks and Mitigations for Agentic AI Security." *OWASP GenAI Security Project*, December 9, 2025. <https://genai.owasp.org/2025/12/09/owasp-genai-security-project-releases-top-10-risks-and-mitigations-for-agentic-ai-security/>, accessed March 2026.

- ⁷⁶ Andrej Karpathy. X.com, February 20, 2026. <https://x.com/karpathy/status/2024987174077432126?s=20>, accessed March 2026.
- ⁷⁷ Isabelle Bousquette, "Here's What OpenClaw Agents are Doing Today." *The Wall Street Journal*, March 18, 2026. <https://www.wsj.com/cio-journal/heres-what-openclaw-agents-are-doing-today-c25d8e6d?st=Qtd5hv>, accessed March 2026.
- ⁷⁸ Isabelle Bousquette, "Here's What OpenClaw Agents are Doing Today." *The Wall Street Journal*, March 18, 2026. <https://www.wsj.com/cio-journal/heres-what-openclaw-agents-are-doing-today-c25d8e6d?st=Qtd5hv>, accessed March 2026.
- ⁷⁹ Lena Hall, Shawn Michaels, and Tarun Chinmai Sekar. "Your OpenClaw Agent Doesn't Sleep. Your Laptop Does. Move It to the Cloud." Akamai Developer Blog, February 9, 2026. <https://www.akamai.com/blog/developers/openclaw-agent-doesnt-sleep-laptop-does-move-cloud>, accessed March 2026.
- ⁸⁰ Lena Hall, Shawn Michaels, and Tarun Chinmai Sekar. "Your OpenClaw Agent Doesn't Sleep. Your Laptop Does. Move It to the Cloud." Akamai Developer Blog, February 9, 2026. <https://www.akamai.com/blog/developers/openclaw-agent-doesnt-sleep-laptop-does-move-cloud>, accessed March 2026.
- ⁸¹ Tom Fosters, "Don't get pinched: the OpenClaw vulnerabilities." *Kaspersky Daily*, February 10, 2026. <https://www.kaspersky.com/blog/openclaw-vulnerabilities-exposed/55263/>, accessed March 2026.
- ⁸² GitHub, "openclaw.ai." *GitHub*. <https://github.com/openclaw/openclaw.ai>, accessed March 2026.
- ⁸³ Jordan Hart, "It's the Mac Mini's moment, thanks to the OpenClaw craze." *Business Insider*, February 19, 2026. <https://www.businessinsider.com/apple-mac-mini-having-a-moment-openclaw-craze-2026-2>, accessed March 2026.
- ⁸⁴ Sebastian Mondragon, "OpenClaw Hit 250K GitHub Stars – Then 20% of Its Skills Were Found Malicious." *Particula*, March 15, 2026. <https://particula.tech/blog/openclaw-security-crisis-malicious-ai-agents>, accessed March 2026.
- ⁸⁵ Deirdre Bosa, "Nvidia CEO Jensen Huang calls OpenClaw 'the most important software release probably ever.'" *CNBC*, March 5, 2026. <https://www.cnb.com/video/2026/03/05/nvidia-ceo-jensen-huang-calls-openclaw-the-most-important-software-release-probably-ever.html>, accessed March 2026.
- ⁸⁶ Ashley Capoot and Ari Levy, "OpenClaw creator Peter Steinberger joining OpenAI, Altman Says." *CNBC*, February 15, 2026. <https://www.cnb.com/2026/02/15/openclaw-creator-peter-steinberger-joining-openai-altman-says.html>, accessed March 2026.
- ⁸⁷ OpenClaw. "What it Does". Openclaw.ai. <https://openclaw.ai/>, accessed March 2026.
- ⁸⁸ The Tech Buzz, "China's OpenClaw Usage Overtakes US in AI Agent Race." *The Tech Buzz*, March 13, 2026, <https://www.techbuzz.ai/articles/china-s-openclaw-usage-overtakes-us-in-ai-agent-race>, accessed March 2026.
- ⁸⁹ Shraddha Goled, "The New Garage: OpenClaw and India's DIY Agent Boom." *Inc42*, March 9, 2026. <https://inc42.com/features/openclaw-indias-diy-ai-agent-boom/>, accessed March 2026.
- ⁹⁰ Salvatore Gariuolo, Vincenzo Ciancaglini, Fernando Tucci, "Viral AI, Invisible Risks: What OpenClaw Reveals About Agentic Assistants." *TrendMicro*, February 6, 2026. https://www.trendmicro.com/en_us/research/26/b/what-openclaw-reveals-about-agentic-assistants.html, accessed March 2026.
- ⁹¹ Louis Columbus, "OpenClaw proves agentic AI works. It also proves your security model doesn't" *Venture Beat*, January 30, 2026. <https://venturebeat.com/security/openclaw-agentic-ai-security-risk-ciso-guide>, accessed March 2026.
- ⁹² Ionut Arghire, "OpenClaw Vulnerability Allowed Websites to Hijack AI Agents." *Security Week*, March 2, 2026. <https://www.securityweek.com/openclaw-vulnerability-allowed-malicious-websites-to-hijack-ai-agents/>, accessed March 2026.
- ⁹³ Alon Klayman, "OpenClaw: The AI Agent Security Crisis Unfolding Right Now." *Reco*, February 12, 2026. <https://www.reco.ai/blog/openclaw-the-ai-agent-security-crisis-unfolding-right-now>, accessed March 2026.
- ⁹⁴ Henry Chandonet, "Meta AI alignment director shares her OpenClaw email-deletion nightmare: 'I had to RUN to my Mac mini'". *Business Insider*, February 23, 2026. <https://www.businessinsider.com/meta-ai-alignment-director-openclaw-email-deletion-2026-2>, accessed March 2026.
- ⁹⁵ Zoe Schiffer and Lauren Goode. "Nvidia is Planning to Launch an Open-Source AI Agent platform." *WIRED*, March 9, 2026. <https://www.wired.com/story/nvidia-planning-ai-agent-platform-launch-open-source/>, accessed March 2026.

⁹⁶ Zoe Schiffer and Lauren Goode. "Nvidia is Planning to Launch an Open-Source AI Agent platform." *WIRED*, March 9, 2026. <https://www.wired.com/story/nvidia-planning-ai-agent-platform-launch-open-source/>, accessed March 2026.

⁹⁷ Zoe Schiffer and Lauren Goode. "Nvidia is Planning to Launch an Open-Source AI Agent platform." *WIRED*, March 9, 2026. <https://www.wired.com/story/nvidia-planning-ai-agent-platform-launch-open-source/>, accessed March 2026.

⁹⁸ Zoe Schiffer and Lauren Goode. "Nvidia is Planning to Launch an Open-Source AI Agent platform." *WIRED*, March 9, 2026. <https://www.wired.com/story/nvidia-planning-ai-agent-platform-launch-open-source/>, accessed March 2026.

⁹⁹ NVIDIA. "NVIDIA Debuts Nemotron 3 Family of Open Models." *NVIDIA Press Release*, December 15, 2025. <https://investor.nvidia.com/news/press-release-details/2025/NVIDIA-Débuts-Nemotron-3-Family-of-Open-Models/default.aspx>, accessed March 2026.

¹⁰⁰ NVIDIA. "NVIDIA Debuts Nemotron 3 Family of Open Models." *NVIDIA Press Release*, December 15, 2025. <https://investor.nvidia.com/news/press-release-details/2025/NVIDIA-Débuts-Nemotron-3-Family-of-Open-Models/default.aspx>, accessed March 2026.

¹⁰¹ Mirko Zorz, "AI went from assistant to autonomous actor and security never caught up." *Help Net Security*, March 3, 2026. <https://www.helpnetsecurity.com/2026/03/03/enterprise-ai-agent-security-2026/>, accessed March 2026.

¹⁰² Zoe Schiffer and Lauren Goode. "Nvidia is Planning to Launch an Open-Source AI Agent platform." *WIRED*, March 9, 2026. <https://www.wired.com/story/nvidia-planning-ai-agent-platform-launch-open-source/>, accessed March 2026.

¹⁰³ Zoe Schiffer and Lauren Goode. "Nvidia is Planning to Launch an Open-Source AI Agent platform." *WIRED*, March 9, 2026. <https://www.wired.com/story/nvidia-planning-ai-agent-platform-launch-open-source/>, accessed March 2026.